

Transfer Learning with MobileNetV2

February 17, 2024

1 Transfer Learning with MobileNetV2

```
[1]: # Note

# A pre-trained model is a network that's already been trained on a large
# ↪ dataset and saved.
# The pre-trained model (MobileNetV2) that we'll be using here has been
# ↪ pre-trained on ImageNet, a dataset
# containing over 14 million images and 1000 classes.
```

```
[2]: ### v1.1
```

```
[3]: # Importing the necessary packages

import matplotlib.pyplot as plt
import numpy as np
import os
import tensorflow as tf
import tensorflow.keras.layers as tfl

from tensorflow.keras.preprocessing import image_dataset_from_directory
from tensorflow.keras.layers.experimental.preprocessing import RandomFlip,
# ↪ RandomRotation
```

1.0.1 Creating the dataset

```
[4]: BATCH_SIZE = 32
IMG_SIZE = (160, 160)
directory = "dataset/"
train_dataset = image_dataset_from_directory(directory,
                                             shuffle=True,
                                             batch_size=BATCH_SIZE,
                                             image_size=IMG_SIZE,
                                             validation_split=0.2,
                                             subset='training',
```

```

seed=42)
validation_dataset = image_dataset_from_directory(directory,
                                                shuffle=True,
                                                batch_size=BATCH_SIZE,
                                                image_size=IMG_SIZE,
                                                validation_split=0.2,
                                                subset='validation',
                                                seed=42)

# The original dataset has some mislabelled images in it as well.

# Note 1

# 'image_size' means all images from the disk will be resized to 160 x 160
→ pixels before being used in training or validation
# datasets.
# The 'image_size' of train_dataset is same as 'image_size' of
→ validation_dataset which is same 'image_size' of test/real world
# dataset on which the model runs. This is because CNNs require a fixed input
→ size and also whatever the model is being
# trained on, it should be validated and tested on similar kind of data.

# shuffle = True ensures that if the minibatches created from the dataset have
→ indices 0,1,2,3,4,...n, then during first epoch,
# the order of minibatches maybe 0,1,2,3....n, but during the second epoch,
→ this order may change.
# Similarly during validation also this order may change.

# setting 'seed=42' ensures that same images are assigned to the training set
→ and the same images are assigned to the
# validation set each time we run the loading_dataset code, given the same
→ dataset and split ratio.

# For seed = x, the training set consists of a set A of images and the
→ validation set consists of a set B of images.
# A and B are non overlapping sets. For seed = y, the training set consists of
→ a set A' of images and the validation set
# consists of a set B' of images. A' and B' are non overlapping sets as well.

# The same 'validation_split' value must be set for both 'train_dataset' and
→ 'validation_dataset'.
# If we set validation_split=0.2, it means that 20% of the dataset will be used
→ for validation, and the remaining 80% will be
# used for training.

# Note 2

```

```

# When we call image_dataset_from_directory and provide the path to the
↳ top-level directory ("dataset/" in this example),
# the function automatically scans the subdirectories (dogs and cats) and
↳ assigns labels to the images based on these
# subdirectory names. In this case, all images in the dogs directory are
↳ labeled as "dogs", and all images in the cats
# directory are labeled as "cats".

# This method of organizing and labeling data is very common because it
↳ simplifies the process of preparing datasets for
# training machine learning models.

```

Found 327 files belonging to 2 classes.

Using 262 files for training.

Found 327 files belonging to 2 classes.

Using 65 files for validation.

```

[5]: # The MobileNetV2 and other pre-trained models are very flexible and accept the
↳ input images in different sizes and not just in
# the size (160,160). This is because they have the GlobalAveragePooling2D
↳ layers before the fully connected layers.

# The GlobalAveragePooling2D layers convert each layer into a single value
↳ making the output size of the pooling layer
# independent of the input image's dimensions. This allows the model to adapt
↳ to different input sizes and feed the
# result into fully connected layers, which do require fixed-size inputs.

```

```

[6]: class_names = train_dataset.class_names

plt.figure(figsize=(10, 10))
for images, labels in train_dataset.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(class_names[labels[i]])
        plt.axis("off")

# When we specify .take(1), it means we are creating a new dataset that
↳ contains only the first mini-batch of our
# training dataset.

```



1.0.2 Preprocessing and Augmentation

```
[7]: AUTOTUNE = tf.data.experimental.AUTOTUNE
train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)

# Using prefetch() prevents a memory bottleneck that can occur when reading
#   ↳ from disk. It sets aside some data and
# keeps it ready for when it's needed, by creating a source dataset from your
#   ↳ input data, applying a transformation to
# preprocess it, then iterating over the dataset one element at a time. Because
#   ↳ the iteration is streaming,
```

```

# the data doesn't need to fit into memory.

# Initial Loading: A batch of data is loaded into memory and processed (e.g.,
↳ decoded, resized, augmented).
# Training: The model begins training on the currently loaded batch.
# Prefetching: While the model is training, the next batch(es) of data are
↳ loaded and prepared in the background.
# The buffer_size determines how many of these batches are prepared in advance.

```

```

[8]: def data_augmenter():
    '''
    Create a Sequential model composed of 2 layers
    Returns:
        tf.keras.Sequential
    '''
    data_augmentation = tf.keras.Sequential()
    data_augmentation.add(RandomFlip('horizontal'))
    data_augmentation.add(RandomRotation(0.2))

    return data_augmentation

# Note

# We can directly pass a list of layers to the tf.keras.Sequential()
↳ constructor.
# Alternatively, we can start with an empty tf.keras.Sequential() model and
↳ then add layers one by one using the .add() method.

# RandomFlip('horizontal') - Flip along the horizontal axis.
# The RandomFlip('horizontal') layer in TensorFlow's Keras API will either flip
↳ an image along the horizontal axis or
# leave it unchanged

# RandomRotation(0.2) - Each image will be rotated by a random angle within the
↳ range of -72 to +72 degrees (0.2*360 = 72)

# Data augmentation doesn't directly increase the number of images in our
↳ dataset.
# This means that each epoch can see slightly different versions of each image,
↳ thereby effectively increasing the diversity of
# the training data without physically increasing the dataset size.
# This also means that the weights see different versions of the same image and
↳ still classify them correctly.

```

```
[9]: augmenter = data_augmenter()

assert(augmenter.layers[0].name.startswith('random_flip')), "First layer must
↳be RandomFlip"
assert augmenter.layers[0].mode == 'horizontal', "RadomFlip parameter must be
↳horizontal"
assert(augmenter.layers[1].name.startswith('random_rotation')), "Second layer
↳must be RandomRotation"
assert augmenter.layers[1].factor == 0.2, "Rotation factor must be 0.2"
assert len(augmenter.layers) == 2, "The model must have only 2 layers"

print('\033[92mAll tests passed!')
```

All tests passed!

```
[10]: data_augmentation = data_augmenter()

for image, _ in train_dataset.take(1):
    plt.figure(figsize=(10, 10))
    first_image = image[0]
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        augmented_image = data_augmentation(tf.expand_dims(first_image, 0))
        plt.imshow(augmented_image[0] / 255)
        plt.axis('off')

# train_dataset.take(1) takes the first minibatch. first_image is the first
↳image in the first minibatch

# tf.expand_dims(first_image, 0) adds an extra dimension to first_image,
↳converting it from a 3D tensor of
# shape [height, width, channels] to a 4D tensor of shape [1, height, width,
↳channels]. This is necessary because
# the data augmentation layers expect a batch of images as input, even if the
↳batch contains only one image.

# The same image undergoes 2 random transformations 9 times and as a result we
↳get 9 variations of the same image.
```




```
[11]: # Points to remember

# When calling image_data_set_from_directory(), specify the train/val subsets,
    ↳ and match the seeds to prevent overlap
# Use prefetch() to prevent memory bottlenecks when reading from disk
# Give your model more to learn from with simple data augmentations like,
    ↳ rotation and flipping.
# When using a pretrained model, it's best to reuse the weights it was trained
    ↳ on.
```

```
[12]: preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input
```

```
# This is the input preprocessing layer of the pre-trained mobilenet_v2 model.
↳ Before we use the pre-trained model on our data,
# it is important to preprocess the data using the same preprocessing layer of
↳ the pre-trained mobilenet_v2 model.
# This is because the weights of the pre-trained model have been trained only
↳ to process certain kinds of input data and
# it is essential for our data to be in that format.
```

1.0.3 Using MobileNetV2 for Transfer Learning

```
[13]: # MobileNetV2 was trained on ImageNet and is optimized to run on mobile and
↳ other low-power applications.
# It's 155 layers deep and very efficient for object detection and image
↳ segmentation tasks, as well as classification tasks.

# Traditional convolutions are often very resource-intensive, and depthwise
↳ separable convolutions are able to reduce the
# number of trainable parameters and operations and also speed up convolutions.
```

```
[14]: IMG_SHAPE = IMG_SIZE + (3,)
# We know IMG_SIZE = (160,160). The tuples are being added.
# Resultant IMG_SHAPE = (160,160,3)
base_model = tf.keras.applications.MobileNetV2(input_shape=IMG_SHAPE,
                                                include_top=True,
                                                weights='imagenet')

# By weights = 'imagenet', it means the weights of MobileNetV2 which have been
↳ trained on the imagenet data.
# include_top=True means the fully-connected output layers that were used for
↳ classification on the
# original dataset (in this case, ImageNet) will be present.
# This is useful when you want to use the model for classification tasks
↳ involving the same number of classes as the
# original dataset (1000 classes for ImageNet).
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_160.h5
14540800/14536120 [=====] - 0s 0us/step

```
[ ]: base_model.summary()
```

```
[16]: # Note the last 2 layers here. They are the so called top layers, and they are
↳ responsible of the classification in the model.

nb_layers = len(base_model.layers)
```



```
print(base_model.layers[nb_layers - 2].name)
print(base_model.layers[nb_layers - 1].name)
```

```
global_average_pooling2d
predictions
```

```
[17]: image_batch, label_batch = next(iter(train_dataset))
      # next(iter(train_dataset)) retrieves the next minibatch

      feature_batch = base_model(image_batch)
      print(feature_batch.shape)
```

```
(32, 1000)
```

```
[18]: #Shows the different label probabilities in one tensor
      label_batch

      # label_batch contains the labels corresponding to each image in image_batch.
      ↳The shape of label_batch is
      # typically (batch_size,), with each entry being the label index for each image.
```

```
[18]: <tf.Tensor: shape=(32,), dtype=int32, numpy=
      array([1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 1, 0,
             0, 0, 1, 0, 1, 1, 1, 1, 0, 0], dtype=int32)>
```

```
[19]: base_model.trainable = False
      # The above line of code means that we are not further training the pre-trained
      ↳weights based on our data

      image_var = tf.Variable(preprocess_input(image_batch))
      # The preprocess_input doesn't automatically convert the image_batch into a
      ↳tensor. We have to do it manually.
      # The above step is not necessary. It isn't necessary to convert into a tensor.
      ↳The base_model will automatically take care of
      # it.
      pred = base_model(image_var)

      tf.keras.applications.mobilenet_v2.decode_predictions(pred.numpy(), top=2)

      # The predictions returned by the base model below follow this format:

      # First the class number, then a human-readable label, and last the probability
      ↳of the image belonging to that class.
      # There are two of these returned for each image in the batch - these the top
      ↳two probabilities returned for that image.
```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json

40960/35363 [=====] - 0s 0us/step

```
[19]: [[('n02489166', 'proboscis_monkey', 0.10329965),
        ('n02102177', 'Welsh_springer_spaniel', 0.07883611)],
       [('n02125311', 'cougar', 0.1654676), ('n02389026', 'sorrel', 0.10764261)],
       [('n02437312', 'Arabian_camel', 0.2923283),
        ('n02437616', 'llama', 0.27713484)],
       [('n03944341', 'pinwheel', 0.31154886), ('n03047690', 'clog', 0.052500293)],
       [('n02454379', 'armadillo', 0.73107153),
        ('n01990800', 'isopod', 0.038719974)],
       [('n02437312', 'Arabian_camel', 0.25663644),
        ('n02422106', 'hartebeest', 0.12122728)],
       [('n02437616', 'llama', 0.6612557),
        ('n02090721', 'Irish_wolfhound', 0.23782855)],
       [('n02133161', 'American_black_bear', 0.82735676),
        ('n02134418', 'sloth_bear', 0.02925945)],
       [('n01518878', 'ostrich', 0.9267562),
        ('n02002724', 'black_stork', 0.0017766367)],
       [('n01518878', 'ostrich', 0.94954586),
        ('n02018795', 'bustard', 0.0028661634)],
       [('n02437616', 'llama', 0.8699833), ('n02412080', 'ram', 0.076757126)],
       [('n02415577', 'bighorn', 0.2429446), ('n02412080', 'ram', 0.160565)],
       [('n02437616', 'llama', 0.9473245), ('n02480495', 'orangutan', 0.0076571796)],
       [('n09428293', 'seashore', 0.48092392), ('n09421951', 'sandbar', 0.26179993)],
       [('n02437312', 'Arabian_camel', 0.95963204),
        ('n02504458', 'African_elephant', 0.0009881927)],
       [('n02509815', 'lesser_panda', 0.9096807),
        ('n02443114', 'polecat', 0.014759211)],
       [('n01518878', 'ostrich', 0.74165), ('n02002724', 'black_stork', 0.07205889)],
       [('n02437312', 'Arabian_camel', 0.49920738),
        ('n02412080', 'ram', 0.11842591)],
       [('n01518878', 'ostrich', 0.87967354),
        ('n02018795', 'bustard', 0.0077298395)],
       [('n02437616', 'llama', 0.82569915),
        ('n02437312', 'Arabian_camel', 0.010480011)],
       [('n01518878', 'ostrich', 0.9612779), ('n02410509', 'bison', 0.0013086519)],
       [('n02437616', 'llama', 0.636178), ('n02412080', 'ram', 0.058401026)],
       [('n02437616', 'llama', 0.5928003), ('n02417914', 'ibex', 0.039721698)],
       [('n02437616', 'llama', 0.83541703), ('n02104029', 'kuvasz', 0.048998024)],
       [('n03042490', 'cliff_dwelling', 0.3091509),
        ('n04208210', 'shovel', 0.06726616)],
       [('n02093647', 'Bedlington_terrier', 0.4338772),
        ('n02113799', 'standard_poodle', 0.4069308)],
       [('n02133161', 'American_black_bear', 0.97880507),
        ('n02132136', 'brown_bear', 0.0055297976)],
       [('n01518878', 'ostrich', 0.83605814), ('n02018795', 'bustard', 0.004823002)],
       [('n02133161', 'American_black_bear', 0.9362426),
```

```
( 'n02134418', 'sloth_bear', 0.007733786)],
[('n03240683', 'drilling_platform', 0.04555222),
 ('n04146614', 'school_bus', 0.033719867)],
[('n02437616', 'llama', 0.9278842),
 ('n02098286', 'West_Highland_white_terrier', 0.0057286685)],
[('n02437616', 'llama', 0.94477594), ('n02423022', 'gazelle', 0.0054335156)]]
```

```
[20]: # here's a whole lot of labels here, some of them hilariously wrong, but none
      ↪ of them say "alpaca."

      # This is because MobileNet pretrained over ImageNet doesn't have the correct
      ↪ labels for alpacas, so when you use the
      # full model, all you get is a bunch of incorrectly classified images.
```

1.0.4 Freezing of weights

```
[21]: def alpaca_model(image_shape=IMG_SIZE, data_augmentation=data_augmenter()):
      ''' Define a tf.keras model for binary classification out of the
      ↪ MobileNetV2 model
      Arguments:
          image_shape -- Image width and height
          data_augmentation -- data augmentation function
      Returns:
      Returns:
          tf.keras.model
      '''

      input_shape = image_shape + (3,)

      base_model = tf.keras.applications.MobileNetV2(input_shape=input_shape,
      include_top=False, # <==
      ↪ Important!!!!
      weights='imagenet') # From
      ↪ imageNet

      # freeze the base model by making it non trainable
      base_model.trainable = False

      # create the input layer (Same as the imageNetv2 input size)
      inputs = tf.keras.Input(shape=input_shape)

      # apply data augmentation to the inputs
      x = data_augmenter()(inputs)
```

```

# Data augmentation doesn't directly increase the number of images in our
→dataset.
# This means that each epoch can see slightly different versions of each
→image, thereby effectively increasing the
# diversity of the training data without physically increasing the dataset
→size.
# This also means that the weights see different versions of the same image
→and still classify them correctly.

# data preprocessing using the same weights the model was trained on
x = preprocess_input(x)

# set training to False to avoid keeping track of statistics in the batch
→norm layer
x = base_model(x, training=False)
# By setting 'training = False', the batch statistics will not be
→influenced by the test set during inference and will be
# influenced only by the training set.
# The BN layers will use the moving averages of mean and variance that were
→computed during training,
# rather than calculating them anew based on the current batch of test
→data, unaffected by the variability in test data.

# Another reason why batch statistics should be influenced only by training
→set and not by test set is that the
# training set minibatch is large whereas it can be only 1 test case that
→we are predicting and we shouldn't allow such a
# small sized test case to influence the batch statistics. If we allow,
→then a single outlier test case will change the
# batch statistics.

# add the new Binary classification layers
# use global avg pooling to summarize the info in each channel
x = tf1.GlobalAveragePooling2D()(x)
# If the input to GlobalAveragePooling2D layer is a tensor of shape
→(batch_size, height, width, channels), the output shape
# after GlobalAveragePooling2D will be (batch_size, channels)

# Since we used 'GlobalAveragePooling2D', we don't need to flatten.

# include dropout with probability of 0.2 to avoid overfitting
x = tf1.Dropout(rate=0.2)(x)
# Note - Dropout is applied only during the training phase. Even during
→validation/testing, dropout is not applied.

```

```

    # use a prediction layer with one neuron (as a binary classifier only needs
    → one)
    outputs = tf.nn.Dense(units=1, activation='linear')(x)

    model = tf.keras.Model(inputs, outputs)

    return model

```

Create your new model using the `data_augmentation` function defined earlier.

```
[22]: model2 = alpaca_model(IMG_SIZE, data_augmentation)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_160_no_top.h5
 9412608/9406464 [=====] - 0s 0us/step

```
[23]: from test_utils import summary, comparator

alpaca_summary = [['InputLayer', [(None, 160, 160, 3)], 0],
                  ['Sequential', (None, 160, 160, 3), 0],
                  ['TensorFlowOpLayer', [(None, 160, 160, 3)], 0],
                  ['TensorFlowOpLayer', [(None, 160, 160, 3)], 0],
                  ['Functional', (None, 5, 5, 1280), 2257984],
                  ['GlobalAveragePooling2D', (None, 1280), 0],
                  ['Dropout', (None, 1280), 0, 0.2],
                  ['Dense', (None, 1), 1281, 'linear']] #linear is the
→ default activation

comparator(summary(model2), alpaca_summary)

for layer in summary(model2):
    print(layer)

```

All tests passed!

```

['InputLayer', [(None, 160, 160, 3)], 0]
['Sequential', (None, 160, 160, 3), 0]
['TensorFlowOpLayer', [(None, 160, 160, 3)], 0]
['TensorFlowOpLayer', [(None, 160, 160, 3)], 0]
['Functional', (None, 5, 5, 1280), 2257984]
['GlobalAveragePooling2D', (None, 1280), 0]
['Dropout', (None, 1280), 0, 0.2]
['Dense', (None, 1), 1281, 'linear']

```

```
[24]: base_learning_rate = 0.001
model2.compile(optimizer=tf.keras.optimizers.Adam(lr=base_learning_rate),
               loss=tf.keras.losses.BinaryCrossentropy(from_logits=True),

```

```
metrics=['accuracy'])
```

```
# Note
# If 'from_logits' = True, the function will internally apply a sigmoid
↳ activation to the logits before calculating the
# binary crossentropy loss. This is useful and numerically stable when the
↳ model's final layer outputs raw logits,
# which are not constrained to [0, 1].
```

```
[25]: initial_epochs = 5
history = model2.fit(train_dataset, validation_data=validation_dataset,
↳ epochs=initial_epochs)
```

Epoch 1/5

9/9 [=====] - 9s 946ms/step - loss: 0.8131 - accuracy: 0.5000 - val_loss: 0.6789 - val_accuracy: 0.5538

Epoch 2/5

9/9 [=====] - 7s 792ms/step - loss: 0.6457 - accuracy: 0.6069 - val_loss: 0.6061 - val_accuracy: 0.5846

Epoch 3/5

9/9 [=====] - 7s 823ms/step - loss: 0.5597 - accuracy: 0.6947 - val_loss: 0.4900 - val_accuracy: 0.6000

Epoch 4/5

9/9 [=====] - 7s 799ms/step - loss: 0.5279 - accuracy: 0.7099 - val_loss: 0.4787 - val_accuracy: 0.6154

Epoch 5/5

9/9 [=====] - 7s 811ms/step - loss: 0.4603 - accuracy: 0.7252 - val_loss: 0.3820 - val_accuracy: 0.6923

```
[26]: # Note

# Yes, TensorFlow's Keras API provides flexibility in how you can fit (train)
↳ your model, allowing for different types of
# input data.

# Method 1

# model.fit(x=inputs, y=outputs, epochs=100)

# Method 2

# model.fit(train_dataset, epochs=100)

# In both the above cases, there is no validation.

# Method 3
```



```
# model.fit(train_dataset, validation_data=validation_dataset, epochs=100)

# In above case, there is validation
```

```
[27]: # Note

# The validation accuracy can sometime be higher than the training accuracy due
    ↳ to following reasons

# 1) If there is strong regularization like dropout, then the model may not
    ↳ perform well on the training set but may perform
# better on the validation dataset because during validation dropout is not
    ↳ applied allowing the full capacity of the model to
# be used.


# 2) Small Dataset Size: If the training dataset is small, the model might not
    ↳ have enough data to learn effectively,
# leading to lower training accuracy. However, if the validation set, by
    ↳ chance, happens to contain examples that are
# easier for the model to classify, then better validation accuracy.
```

```
[28]: acc = [0.] + history.history['accuracy']
val_acc = [0.] + history.history['val_accuracy']

loss = history.history['loss']
val_loss = history.history['val_loss']

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),1])
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([0,1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()
```



```
[29]: class_names
```

```
[29]: ['alpaca', 'not alpaca']
```

The results are ok, but could be better. Next, try some fine-tuning.

1.0.5 Fine-tuning the Model

```
[30]: # Note

# For fine tuning a pre-trained model, we must use small learning rate to train
# the later layers on our data because the
```

```
# weights have already been learnt and it is quite close to the minima. So if
↳ learning rate is high, the optimizer may diverge.
# We need to use small learning rate because the weights are not learning for
↳ the first time but they are in the final stages of
# learning.
```

```
[31]: # Note

# When the network is in its earlier stages, it trains on low-level features,
↳ like edges. In the later layers,
# more complex, high-level features like wispy hair or pointy ears begin to
↳ emerge. For transfer learning,
# the low-level features can be kept the same, as they have common features for
↳ most images. When you add new data,
# you generally want the high-level features to adapt to it, which is rather
↳ like letting the network learn to
# detect features more related to your data, such as soft fur or big teeth.

# To achieve this, just unfreeze the final layers and re-run the optimizer with
↳ a smaller learning rate,
# while keeping all the other layers frozen.
```

```
[32]: base_model = model2.layers[4]
# To understand the above line of code, see the model2 layers in the code cell
↳ below.
# The MobileNetV2 model is model2.layers[4]
base_model.trainable = True
# The above line of code makes all the layers of MobileNetV2 model trainable
# Let's take a look to see how many layers are in the base model
print("Number of layers in the base model: ", len(base_model.layers))

# Fine-tune from this layer onwards
fine_tune_at = 120

# Freeze all the layers before the `fine_tune_at` layer
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False

# Define a BinaryCrossentropy loss function. Use from_logits=True
loss_function=tf.keras.losses.BinaryCrossentropy(from_logits=True)
# Define an Adam optimizer with a learning rate of 0.1 * base_learning_rate
optimizer = tf.keras.optimizers.Adam(lr=0.1*base_learning_rate)
# Use accuracy as evaluation metric
metrics=['accuracy']
```

```

model2.compile(loss=loss_function,
               optimizer = optimizer,
               metrics=metrics)

# Note

# So even if we set 1 layer as trainable, then it is required to compile the
→ model and define the optimizer,
# the loss function and metrics.

```

Number of layers in the base model: 155

```

[34]: for i, layer in enumerate(model2.layers):
       print(i, layer.name)

```

```

0 input_3
1 sequential_3
2 tf_op_layer_RealDiv
3 tf_op_layer_Sub
4 mobilenetv2_1.00_160
5 global_average_pooling2d_1
6 dropout
7 dense

```

```

[35]: assert type(loss_function) == tf.python.keras.losses.BinaryCrossentropy, "Not
→ the correct layer"
assert loss_function.from_logits, "Use from_logits=True"
assert type(optimizer) == tf.keras.optimizers.Adam, "This is not an Adam
→ optimizer"
assert optimizer.lr == base_learning_rate / 10, "Wrong learning rate"
assert metrics[0] == 'accuracy', "Wrong metric"

print('\033[92mAll tests passed!')

```

All tests passed!

```

[38]: fine_tune_epochs = 5
total_epochs = initial_epochs + fine_tune_epochs

history_fine = model2.fit(train_dataset,
                          epochs=total_epochs,
                          initial_epoch=history.epoch[-1],
                          validation_data=validation_dataset)

# When we specify the 'initial_epoch' parameter, the epoch number starts from
→ 'initial_epoch' + 1 and goes upto
# total_epochs (inclusive)

```

```

Epoch 5/10
9/9 [=====] - 9s 945ms/step - loss: 0.1523 - accuracy:
0.9389 - val_loss: 0.0681 - val_accuracy: 0.9846
Epoch 6/10
9/9 [=====] - 8s 943ms/step - loss: 0.1167 - accuracy:
0.9618 - val_loss: 0.0632 - val_accuracy: 0.9538
Epoch 7/10
9/9 [=====] - 9s 957ms/step - loss: 0.1259 - accuracy:
0.9389 - val_loss: 0.1120 - val_accuracy: 0.9692
Epoch 8/10
9/9 [=====] - 8s 923ms/step - loss: 0.1065 - accuracy:
0.9504 - val_loss: 0.0918 - val_accuracy: 0.9846
Epoch 9/10
9/9 [=====] - 8s 923ms/step - loss: 0.1056 - accuracy:
0.9695 - val_loss: 0.2592 - val_accuracy: 0.8769
Epoch 10/10
9/9 [=====] - 8s 936ms/step - loss: 0.1282 - accuracy:
0.9389 - val_loss: 0.1030 - val_accuracy: 0.9231

```

```
[37]: print(history.epoch)
```

```
[0, 1, 2, 3, 4]
```

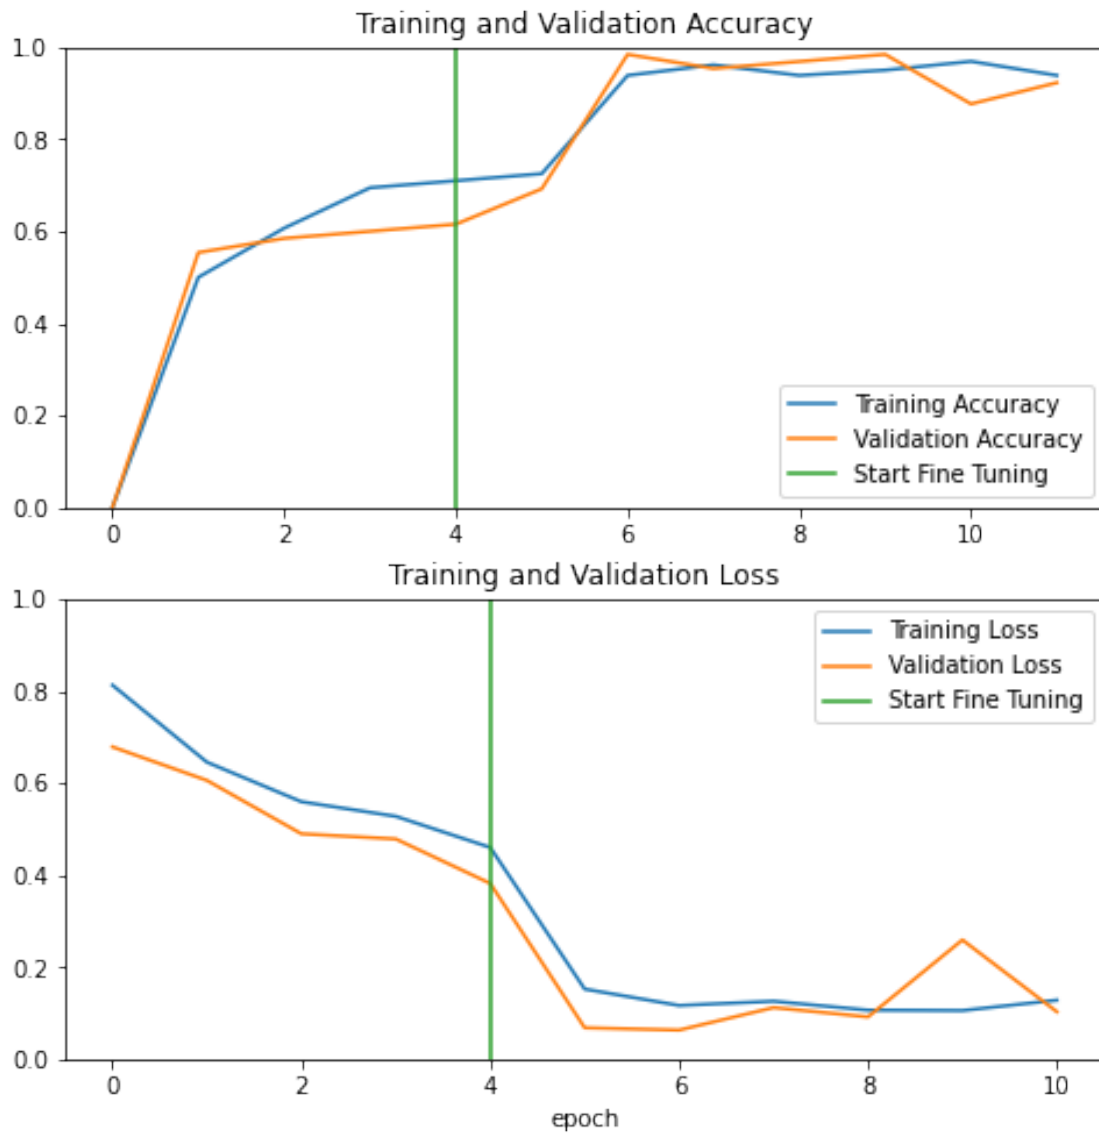
```
[39]: acc += history_fine.history['accuracy']
      val_acc += history_fine.history['val_accuracy']

      loss += history_fine.history['loss']
      val_loss += history_fine.history['val_loss']
```

```
[40]: plt.figure(figsize=(8, 8))
      plt.subplot(2, 1, 1)
      plt.plot(acc, label='Training Accuracy')
      plt.plot(val_acc, label='Validation Accuracy')
      plt.ylim([0, 1])
      plt.plot([initial_epochs-1, initial_epochs-1],
               plt.ylim(), label='Start Fine Tuning')
      plt.legend(loc='lower right')
      plt.title('Training and Validation Accuracy')

      plt.subplot(2, 1, 2)
      plt.plot(loss, label='Training Loss')
      plt.plot(val_loss, label='Validation Loss')
      plt.ylim([0, 1.0])
      plt.plot([initial_epochs-1, initial_epochs-1],
               plt.ylim(), label='Start Fine Tuning')
      plt.legend(loc='upper right')
      plt.title('Training and Validation Loss')
```

```
plt.xlabel('epoch')
plt.show()
```



```
[ ]: # Note

# The entire model was initially trained for 5 epochs with all the weights of MobileNetV2 frozen.
# Then the entire model (already trained for 5 epochs) was trained for 5 more epochs with only the first 120 layer weights of MobileNetV2 frozen and later layer weights trainable.
```



```
[ ]: # Note

# How fine tuning also helps batch normalization layers?

# Earlier when we freezed the BN layer parameters, the batch statistics (moving
  ↳average of mean and variance) were based on the
# imageNet dataset. The batch normalized version of imageNet dataset is of
  ↳distribution A whereas the
# batch normalized version of the test dataset is of distribution B. A and B
  ↳are different. The distribution B inputs are not
# familiar to the BN layer weights and hence these weights don't predict
  ↳efficiently.

# Now we unfreeze the BN layer parameters. Now the saved batch statistics
  ↳(moving average of mean and variance) in memory are
# based on the custom training dataset and also based on previously imageNet
  ↳dataset. if more fine-tuning is done,
# then the batch statistics (moving average of mean and variance) will be more
  ↳influenced by the moving average of custom
# training dataset. The batch normalized version of the custom training dataset
  ↳is of say distribution x.
# The batch normalized version of the test dataset is of say distribution y . x
  ↳and y are pretty close (if proper
# fine-tuning is done). The BN layer weights see a familiar distribution and
  ↳they predict efficiently.

# The new moving averages is found by considering the earlier moving average of
  ↳the imageNet and the moving average of the
# current custom training dataset.
```

```
[ ]: # Note

# To adapt the classifier to new data: Delete the top layer, add a new
  ↳classification layer, and train only on that layer.
# Fine-tune the final layers of your model to capture high-level details near
  ↳the end of the network and
# potentially improve accuracy.
```